

Data Exploration

May 5, 2020

```
[3]: import os
print("cwd:", os.getcwd())
```

cwd: /Users/jsennett/Code/top-k-insights/report/notebooks

```
[209]: import pandas as pd
import numpy as np
import scipy
import heapq
import scipy.stats
import matplotlib.pyplot as plt
```

0.1 Number of papers per author

```
[ ]: filename = "/Users/jsennett/Code/top-k-insights/data/
↳all-author-paper-counts-by-year.csv"
df = pd.read_csv(filename, encoding='mac_roman')
```

```
[30]: print(len(data), "rows (4.4 million)")
print(df.columns)
df.head()
```

4395104 rows (4.4 million)

Index(['num_papers', 'year', 'authid', 'rank', 'log_rank', 'log_num_papers'],
dtype='object')

```
[30]:
```

	num_papers	year	authid	rank	log_rank	log_num_papers
0	1	1990	0	1571776.0	20.583964	0.000000
1	4	1990	1	329751.0	18.331018	2.000000
2	5	1990	3	230373.0	17.813612	2.321928
3	1	1990	4	1571777.0	20.583965	0.000000
4	1	1990	7	1571778.0	20.583966	0.000000


```
[20]: df.describe()
```

```
[20]:
```

	num_papers	year	authid
count	4.395104e+06	4.395104e+06	4.395104e+06
mean	1.980562e+00	2.007571e+03	5.751863e+05
std	2.490828e+00	6.059493e+00	4.682082e+05
min	1.000000e+00	1.990000e+03	0.000000e+00
25%	1.000000e+00	2.004000e+03	1.710480e+05
50%	1.000000e+00	2.009000e+03	4.489020e+05
75%	2.000000e+00	2.012000e+03	9.167660e+05
max	2.000000e+02	2.015000e+03	1.691340e+06


```
[23]: # Scatter plots of some relevant measures
df.plot.scatter('year', 'num_papers')
```

```
[23]: <matplotlib.axes._subplots.AxesSubplot at 0x12a1a0d30>
```



```
[28]: # Plot num_papers versus rank
df['rank'] = df['num_papers'].rank(method='first', ascending=False)
df.plot.scatter('rank', 'num_papers')
```


```
[28]: <matplotlib.axes._subplots.AxesSubplot at 0x127102e80>
```



```
[29]: df['log_rank'] = scipy.log2(df['rank'])
      df['log_num_papers'] = scipy.log2(df['num_papers'])
      df.plot.scatter('log_rank', 'log_num_papers')
```

```
/Users/jsennett/.pyenv/versions/3.6.1/envs/topk/lib/python3.6/site-
packages/ipykernel_launcher.py:1: DeprecationWarning: scipy.log2 is deprecated
and will be removed in SciPy 2.0.0, use numpy.lib.scimath.log2 instead
    """Entry point for launching an IPython kernel.
/Users/jsennett/.pyenv/versions/3.6.1/envs/topk/lib/python3.6/site-
packages/ipykernel_launcher.py:2: DeprecationWarning: scipy.log2 is deprecated
and will be removed in SciPy 2.0.0, use numpy.lib.scimath.log2 instead
```

```
[29]: <matplotlib.axes._subplots.AxesSubplot at 0x12a110eb8>
```



0.2 All papers

```
[5]: papers_filename = "/Users/jsennett/Code/top-k-insights/data/all-papers.csv"
p_df = pd.read_csv(papers_filename, encoding='mac_roman', dtype = {'school':  
    ↳str})
```

```
[6]: print(len(p_df), "rows (4.4 million)")
print(p_df.columns)
p_df.head()
```

2991406 rows (4.4 million)

Index(['paperid', 'venue_name', 'year', 'school', 'venue_type'], dtype='object')


```
[6]:
```

	paperid	venue_name	year	school	venue_type
0	5389	Future Generation Comp. Syst.	2004	NaN	0
1	5390	Future Generation Comp. Syst.	2010	NaN	0
2	5407	Future Generation Comp. Syst.	2009	NaN	0
3	5414	Future Generation Comp. Syst.	2001	NaN	0
4	5449	Future Generation Comp. Syst.	2004	NaN	0

```
[7]: # Insight: positive trend of paper counts over year considering the full set of  
    ↳data
p_df.groupby('year')\
```


```
.agg('count')\
.reset_index()\
.rename(columns={'paperid': 'number of papers'})\
.plot.scatter('year', 'number of papers')
```

[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1035bbdd8>



```
[8]: # Insight: positive trend of paper counts over year considering the full set of
      ↪ data
p_df[p_df['venue_type'] == 1].groupby('year')\
    .agg('count')\
    .reset_index()\
    .rename(columns={'paperid': 'number of papers'})\
    .plot.scatter('year', 'number of papers')
```

[8]: <matplotlib.axes._subplots.AxesSubplot at 0x10412fd30>



```
[9]: p_df[p_df['venue_type'] == 0]['year'].value_counts()
```


```
[9]: 2014    113308
      2015    111441
      2013    107439
      2012    100868
      2011     92914
      2010     86390
      2009     79643
      2008     72474
      2007     66958
      2006     59665
      2005     52474
      2004     44429
      2003     40648
      2002     35316
      2001     32921
      2000     30846
      1999     26084
      1998     23625
      1997     21491
      1996     20078
      1995     18390
      1994     16810
      1993     15366
```

```
1992      13455
1991      11976
1990      10266
Name: year, dtype: int64
```

```
[10]: p_df['M'] = 1
```

```
[11]: # Insight: positive trend of paper counts over year considering the full set of
      ↪ data
p_df[p_df['venue_type'] == 1].groupby('year')\
    .agg('count')\
    .reset_index()\
    .rename(columns={'paperid': 'number of papers'})\
    .plot.scatter('year', 'number of papers')
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1166d72b0>
```




```
[32]: def scatter(df, dimension):
      return df.groupby(dimension)\
          .agg('count')\
          .reset_index()\
          .rename(columns={'paperid': 'number of papers'})\
          .plot.scatter(dimension, 'number of papers')
```

```
[23]: def rank_scatter(df, dimension):
        rank_df = df.groupby(dimension).agg('count').reset_index()
        rank_df['rank'] = rank_df['M'].rank(method='first', ascending=False)
        return rank_df.plot.scatter(x='rank', y='M', label='rank of ' + dimension)
```


```
[26]: rank_scatter(p_df, 'year')
```

```
[26]: <matplotlib.axes._subplots.AxesSubplot at 0x1169f4a90>
```



```
[15]: venue_ranks['rank'] = venue_ranks['M'].rank(method='first', ascending=False)
        venue_ranks.plot.scatter(x='rank', y='M')
```

```
[15]: <matplotlib.axes._subplots.AxesSubplot at 0x116402128>
```

```
[20]: p_df['paperid'].rank()
```

```
[20]: 0          4674.0
      1          4675.0
      2          4692.0
      3          4699.0
      4          4734.0
      ...
      2991401    4583.0
      2991402    4617.0
      2991403    4627.0
      2991404    4646.0
      2991405    4653.0
      Name: paperid, Length: 2991406, dtype: float64
```

```
[ ]: p_df.groupby('venue_name').agg('count').head().rank()

df['rank'] = df['num_papers'].rank(method='first', ascending=False)
```

```
[41]: result_set = p_df.groupby(['venue_name', 'year']).agg({'M': 'sum'}).reset_index()
      result_set['M'] = result_set.sort_values("year")['M'].diff()
      result_set.head()
```

```
[41]:
```

	venue_name	year	M
0	"CloudCom-Asia	2015	7.0
1	#MSM	2011	-3.0
2	#MSM	2012	-13.0
3	#MSM	2013	-74.0
4	#MSM	2014	-78.0

```
[ ]: result_set['absM'] = abs(result_set['M'])
```

```
[ ]: rank_scatter(result_set, 'absM')
```

```
[49]: import scipy
result_set['logabsM'] = scipy.log2(result_set['absM'])
result_set['logrank'] = scipy.log2(result_set['M'].rank(method='first',
↪ascending=False))
```

/Users/jsennett/.pyenv/versions/3.6.1/envs/topk/lib/python3.6/site-packages/ipykernel_launcher.py:2: DeprecationWarning: scipy.log2 is deprecated and will be removed in SciPy 2.0.0, use numpy.lib.scimath.log2 instead

/Users/jsennett/.pyenv/versions/3.6.1/envs/topk/lib/python3.6/site-packages/numpy/lib/scimath.py:122: RuntimeWarning: invalid value encountered in less

```
if any(isreal(x) & (x < 0)):
```

/Users/jsennett/.pyenv/versions/3.6.1/envs/topk/lib/python3.6/site-packages/numpy/lib/scimath.py:412: RuntimeWarning: divide by zero encountered in log2


```
return nx.log2(x)
```

/Users/jsennett/.pyenv/versions/3.6.1/envs/topk/lib/python3.6/site-packages/ipykernel_launcher.py:3: DeprecationWarning: scipy.log2 is deprecated and will be removed in SciPy 2.0.0, use numpy.lib.scimath.log2 instead

This is separate from the ipykernel package so we can avoid doing imports until


```
[63]: result_set['rank'] = result_set['M'].rank(method='first', ascending=False)
result_set.plot.scatter('rank', 'M')
```

```
[63]: <matplotlib.axes._subplots.AxesSubplot at 0x12d765ba8>
```



```
[68]: result_set[result_set['M'] > 0].plot.scatter('logrank', 'logabsM')
```

```
[68]: <matplotlib.axes._subplots.AxesSubplot at 0x1242ac358>
```



```
[ ]: result_set[result_set['M'] > 50].plot.scatter('rank', 'M')
```

```
[ ]: cutoff = 0.2
result_set[result_set['rank'] < cutoff * len(result_set)].plot.
↳scatter('logrank', 'logabsM')
```

```
[108]: # pct of total:
pct_of_total = p_df.copy()
```

```
[109]: pct_of_total = pct_of_total.groupby(['venue_name']).agg({'M': 'sum'})
pct_of_total = pct_of_total.reset_index(drop=False)
pct_of_total.head()
```

```
[109]:
```

	venue_name	M
0	"CloudCom-Asia	29
1	#MSM	64
2	10th Anniversary Colloquium of UNU/IIST	26
3	15. WLP	12
4	1999 ACM SIGMOD Workshop on Research Issues in...	12


```
[110]: pct_of_total['M'] = 100 * pct_of_total['M'] / sum(pct_of_total['M'])
```

```
[ ]: pct_of_total['rank'] = pct_of_total['M'].rank(method='first', ascending=False)
pct_of_total.plot.scatter('rank', 'M')
```

```
[ ]: pct_of_total['logrank'] = scipy.log2(pct_of_total['rank'])
pct_of_total['logM'] = scipy.log2(pct_of_total['M'])
pct_of_total.plot.scatter('logrank', 'logM')
```

```
[114]: pct_of_total.plot.scatter('logrank', 'M')
```

```
[114]: <matplotlib.axes._subplots.AxesSubplot at 0x11e80d208>
```



```
[ ]: pct_of_total['M'] = 100 * result_set['M'] / sum(result_set['M'])
result_set['M'] = 100 * result_set[self.measure] / sum(result_set[self.measure])
p_df.head()
```

```
[127]: # Validate composite extractor result sets
"""
INFO:root:Valid SG/CE combo: subspace({'year': 2000}), dim(venue_type),
↳ce(['count', 'count'], ['pct', 'year'])
INFO:root:extract_result_set({'year': 2000}, venue_type, [['count', 'count'],
↳['pct', 'year'])
INFO:root:RESULT SET:
      venue_type  year  count  mod2      M
10             0  2000  30846    *  2.363180
36             1  2000  36539    *  2.175786
62             3  2000    43     *  0.633844
"""
print(100 * len(p_df[(p_df['year']==2000) & (p_df['venue_type'] == 0)]) /
↳len(p_df[p_df['venue_type']==0]))
print(100 * len(p_df[(p_df['year']==2000) & (p_df['venue_type'] == 1)]) /
↳len(p_df[p_df['venue_type']==1]))
print(100 * len(p_df[(p_df['year']==2000) & (p_df['venue_type'] == 3)]) /
↳len(p_df[p_df['venue_type']==3]))
```

2.3631801727605293

2.175786183558252

0.6338443396226415


```
[ ]: """
INFO:root:extract_result_set({'year': 2000}, venue_type, [['count', 'count'],
↳['delta_prev', 'year']]
returning result set:

    venue_type  year  count mod2    M
1           0  2000  30846    *  4762
3           1  2000  36539    *  3805
5           3  2000     43    *     3
"""

print(len(p_df[(p_df['year'] == 2000) & (p_df['venue_type'] == 0)]) -
↳len(p_df[(p_df['year'] == 1999) & (p_df['venue_type'] == 0)]))
print(len(p_df[(p_df['year'] == 2000) & (p_df['venue_type'] == 1)]) -
↳len(p_df[(p_df['year'] == 1999) & (p_df['venue_type'] == 1)]))
print(len(p_df[(p_df['year'] == 2000) & (p_df['venue_type'] == 3)]) -
↳len(p_df[(p_df['year'] == 1999) & (p_df['venue_type'] == 3)]))
```

```
[134]: x = pd.DataFrame([10 * np.random.normal() + 15 for _ in range(100)],
↳columns=['M'])
x['M'].hist()
```

[134]: <matplotlib.axes._subplots.AxesSubplot at 0x11efa0ba8>



```
[ ]: max(x['M'])
      scipy.stats.norm.pdf()
```


```
[ ]: (38 - x['M'].mean()) / x['M'].std()
```

```
[ ]: scipy.stats.norm.ppf(100/)
```

```
[ ]: x.rank(method='first', ascending=False)
```

```
[147]: x.hist()
```

```
[147]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x13f067eb8>]],
      dtype=object)
```



```
[156]: x['rank'] = x['M'].rank(ascending=True, method='first')
      x['pct'] = x['rank'] / (len(x)+1)
      x['Z'] = (x['M'] - x['M'].mean()) / x['M'].std()
      x['predZ'] = scipy.stats.norm.pdf(x['pct'])
      x['pred'] = x['pct'] * x['M'].std() + x['M'].mean()
```

```
[ ]: x['pred'].hist()
```

```
[166]: x.head(25)
```

```
[166]:
```

	M	rank	pct	Z	predZ	pred
0	6.771423	22.0	0.217822	-0.700520	0.389589	15.949612
1	38.250489	100.0	0.990099	2.449180	0.244366	23.667987
2	-8.598282	3.0	0.029703	-2.238366	0.398766	14.069496
3	19.697190	73.0	0.722772	0.592793	0.307236	20.996242
4	-2.258100	6.0	0.059406	-1.603987	0.398239	14.366356
5	28.667375	94.0	0.930693	1.490323	0.258714	23.074266
6	7.262322	25.0	0.247525	-0.651402	0.386906	16.246473
7	30.158047	96.0	0.950495	1.639475	0.253940	23.272173
8	24.714445	85.0	0.841584	1.094804	0.279971	22.183684
9	-0.213637	8.0	0.079208	-1.399424	0.397693	14.564263
10	16.941651	62.0	0.613861	0.317082	0.330433	19.907753
11	-11.494205	1.0	0.009901	-2.528124	0.398923	13.871589
12	13.252900	47.0	0.465347	-0.052003	0.358004	18.423450
13	13.814717	52.0	0.514851	0.004211	0.349422	18.918218
14	7.703141	27.0	0.267327	-0.607295	0.384939	16.444380
15	2.012392	13.0	0.128713	-1.176694	0.395651	15.059031
16	21.757046	76.0	0.752475	0.798896	0.300578	21.293102
17	23.147003	80.0	0.792079	0.937971	0.291524	21.688916
18	20.522275	74.0	0.732673	0.675349	0.305030	21.095195
19	24.218423	84.0	0.831683	1.045174	0.282299	22.084730
20	17.951269	65.0	0.643564	0.418101	0.324320	20.204614
21	7.833889	28.0	0.277228	-0.594213	0.383903	16.543334
22	23.822498	82.0	0.811881	1.005559	0.286931	21.886823
23	23.863100	83.0	0.821782	1.009621	0.284620	21.985777
24	11.288491	42.0	0.415842	-0.248556	0.365898	17.928683

```
[176]: x.head()
```

```
[176]:
```

	M
0	6.771423
1	38.250489
2	-8.598282
3	19.697190
4	-2.258100

```
[177]: xmean = x['M'].mean()
xstd = x["M"].std()
```

```
[ ]: x['rank'] = x['M'].rank(method='first')
x.plot.scatter('rank', 'M')
```

```
[184]:
```

```
[193]: (1 - scipy.stats.norm.cdf(x_max_Z)) / (.01)
```

```
[193]: 0.5536938553650228
```



```
[198]: alpha = 1/100
x_max_p = 1 - scipy.stats.norm.cdf(x_max_Z)
print(x_max_p)
```

0.005536938553650228


```
[ ]: significance = max(1 - x_max_p / alpha, 0.0)
print(significance)
```

```
[ ]:
```

```
[234]: x = pd.DataFrame([10 * np.random.normal() + 15 for _ in range(100)],
↳ columns=['M'])
x_max = max(x['M'])
x_mean = x.drop(x['M'].idxmax())['M'].mean()
x_std = x.drop(x['M'].idxmax())['M'].std()
x_max_Z = (x_max - x_mean)/x_std
x_max_p = 1 - scipy.stats.norm.cdf(x_max_Z)

x.hist()
plt.axvline(scipy.stats.norm.ppf(.99) * x_std + x_mean, color='r')
plt.axvline(x_max, color='g')
print("significance:", max(1 - x_max_p / alpha, 0.0))
```

significance: 0.0



```
[ ]: x = pd.DataFrame([10 * np.random.normal() + 15 for _ in range(100)],
    ↪columns=['M'])
x_max = max(x['M'])
x_mean = x.drop(x['M'].idxmax())['M'].mean()
x_std = x.drop(x['M'].idxmax())['M'].std()
x_max_Z = (x_max - x_mean)/x_std
x_max_p = 1 - scipy.stats.norm.cdf(x_max_Z)

x.hist()
plt.axvline(scipy.stats.norm.ppf(.99) * x_std + x_mean, color='r')
plt.axvline(x_max, color='g')
print("significance:", max(1 - x_max_p / alpha, 0.0))
```

```
[346]: def simulate(n):
    x = pd.DataFrame([10 * np.random.normal() + 15 for _ in range(n)],
    ↪columns=['M'])
    x_max = max(x['M'])
    alpha = 1/n

    x_mean = x['M'].mean()
    x_std = x['M'].std()
    x_max_Z = (x_max - x_mean)/x_std
    x_max_p = 1 - scipy.stats.norm.cdf(x_max_Z)


    x.hist(density=True)
    plt.axvline(scipy.stats.norm.ppf(1-1/n) * x_std + x_mean, color='r',
    ↪label='expected max val')
    plt.axvline(x_max, color='g', label='actual max val')

    # plot normal dist
    x_axis = np.arange(min(x['M']), max(x['M']), .001)
    plt.plot(x_axis, scipy.stats.norm.pdf(x_axis, x_mean, x_std))
    # Mean = 0, SD = 2.

    print("significance:", max(1 - x_max_p / alpha, 0.0))
    print("x_max_p", x_max_p)
    print("alpha (1/n)", alpha)
```


```
[348]: simulate(1000)
```

```
significance: 0.9971802493782258
x_max_p 2.8197506217741974e-06
alpha (1/n) 0.001
```




```
[351]: simulate(1000)
```

```
significance: 0.0  
x_max_p 0.0014602342050724637  
alpha (1/n) 0.001
```




```
[353]: simulate(1000)
```

```
significance: 0.35295361397935565  
x_max_p 0.0006470463860206443  
alpha (1/n) 0.001
```




```
[241]: simulate(1000)
```

significance: 0.997779697503486



```
[279]: simulate(1000)
```

significance: 0.9470496492825933



```
[ ]: scipy.stats.norm.ppf(.5, x_mean, x_std)
```


```
[ ]: scipy.stats.norm.pdf(x_mean, x_mean, x_std)
```

```
[290]: x_std
```

```
[290]: 9.039491770341439
```

```
[293]: x_min = -25
x_max = 50
mean = 15
std = 9
x = np.linspace(x_min, x_max, 100)
y = scipy.stats.norm.pdf(x, mean, std) * 1000
plt.plot(x, y)
```

```
[293]: [<matplotlib.lines.Line2D at 0x131f88dd8>]
```



```
[ ]: import numpy as np
      from scipy.stats import norm
      import matplotlib.pyplot as plt
```

```
[ ]: # Plot the histogram.
plt.hist(data, bins=25, density=True, alpha=0.6, color='g')

# Plot the PDF.
xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
p = norm.pdf(x, mu, std)
plt.plot(x, p, 'k', linewidth=2)
title = "Fit results: mu = %.2f, std = %.2f" % (mu, std)
plt.title(title)

plt.show()
```

```
[ ]: p_df[p_df['venue_name'] == 'CoRR']
```

```
[ ]: result_set = p_df
      analysis_dimension = 'venue_name'
      subspace = {'venue_name': 'CoRR'}
      result_set[result_set[analysis_dimension] == subspace[analysis_dimension]]
```

```
[405]: p_df = p_df.fillna('')
school_size = p_df.groupby('school').agg({"M": "sum"})
school_size[school_size['M'] > len(p_df) * .01]
```

```
[405]:          M
school
2984622
```

```
[ ]: p_df = pd.read_csv(papers_filename, encoding='mac_roman', dtype = {'school':  
↳str})
```

```
[409]: len(p_df)
```

```
[409]: 0.9493745166098949
```